

## 1) Load CIFAR-100 dataset and Apply Convolutional Neural Network (CNN) to train and compile the model.

```

1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import tensorflow as tf
5 import keras
6 import matplotlib.pyplot as plt
7 import math
8 import datetime
9 import platform
10
11
12

```

```
1 tf.keras.datasets.cifar10.load_data()
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>

170498071/170498071 [=====] - 4s 0us/step

```

((array([[[[ 59,  62,  63],
            [ 43,  46,  45],
            [ 50,  48,  43],
            ...,
            [158, 132, 108],
            [152, 125, 102],
            [148, 124, 103]],
          [[ 16,  20,  20],
            [  0,   0,   0],
            [ 18,   8,   0],
            ...,
            [123,  88,  55],
            [119,  83,  50],
            [122,  87,  57]],
          [[ 25,  24,  21],
            [ 16,   7,   0],
            [ 49,  27,   8],
            ...,
            [118,  84,  50],
            [120,  84,  50],
            [109,  73,  42]],
          ...,
          [[208, 170,  96],
            [201, 153,  34],
            [198, 161,  26],
            ...,
            [160, 133,  70],
            [ 56,  31,   7],
            [ 53,  34,  20]],
          [[180, 139,  96],
            [173, 123,  42],
            [186, 144,  30],
            ...,
            [184, 148,  94],
            [ 97,  62,  34],
            [ 83,  53,  34]],
          [[177, 144, 116],
            [168, 129,  94],
            [179, 142,  87],
            ...,
            [216, 184, 140],
            [151, 118,  84],
            [123,  92,  72]]],
      [[[154, 177, 187],
        [126, 137, 136],
        [105, 104,  95],
        ...,
        [ 91,  95,  71],

```

```
1 (train_images, train_labels), (test_images, test_labels) = keras.datasets.cifar10.load_data()
```

```

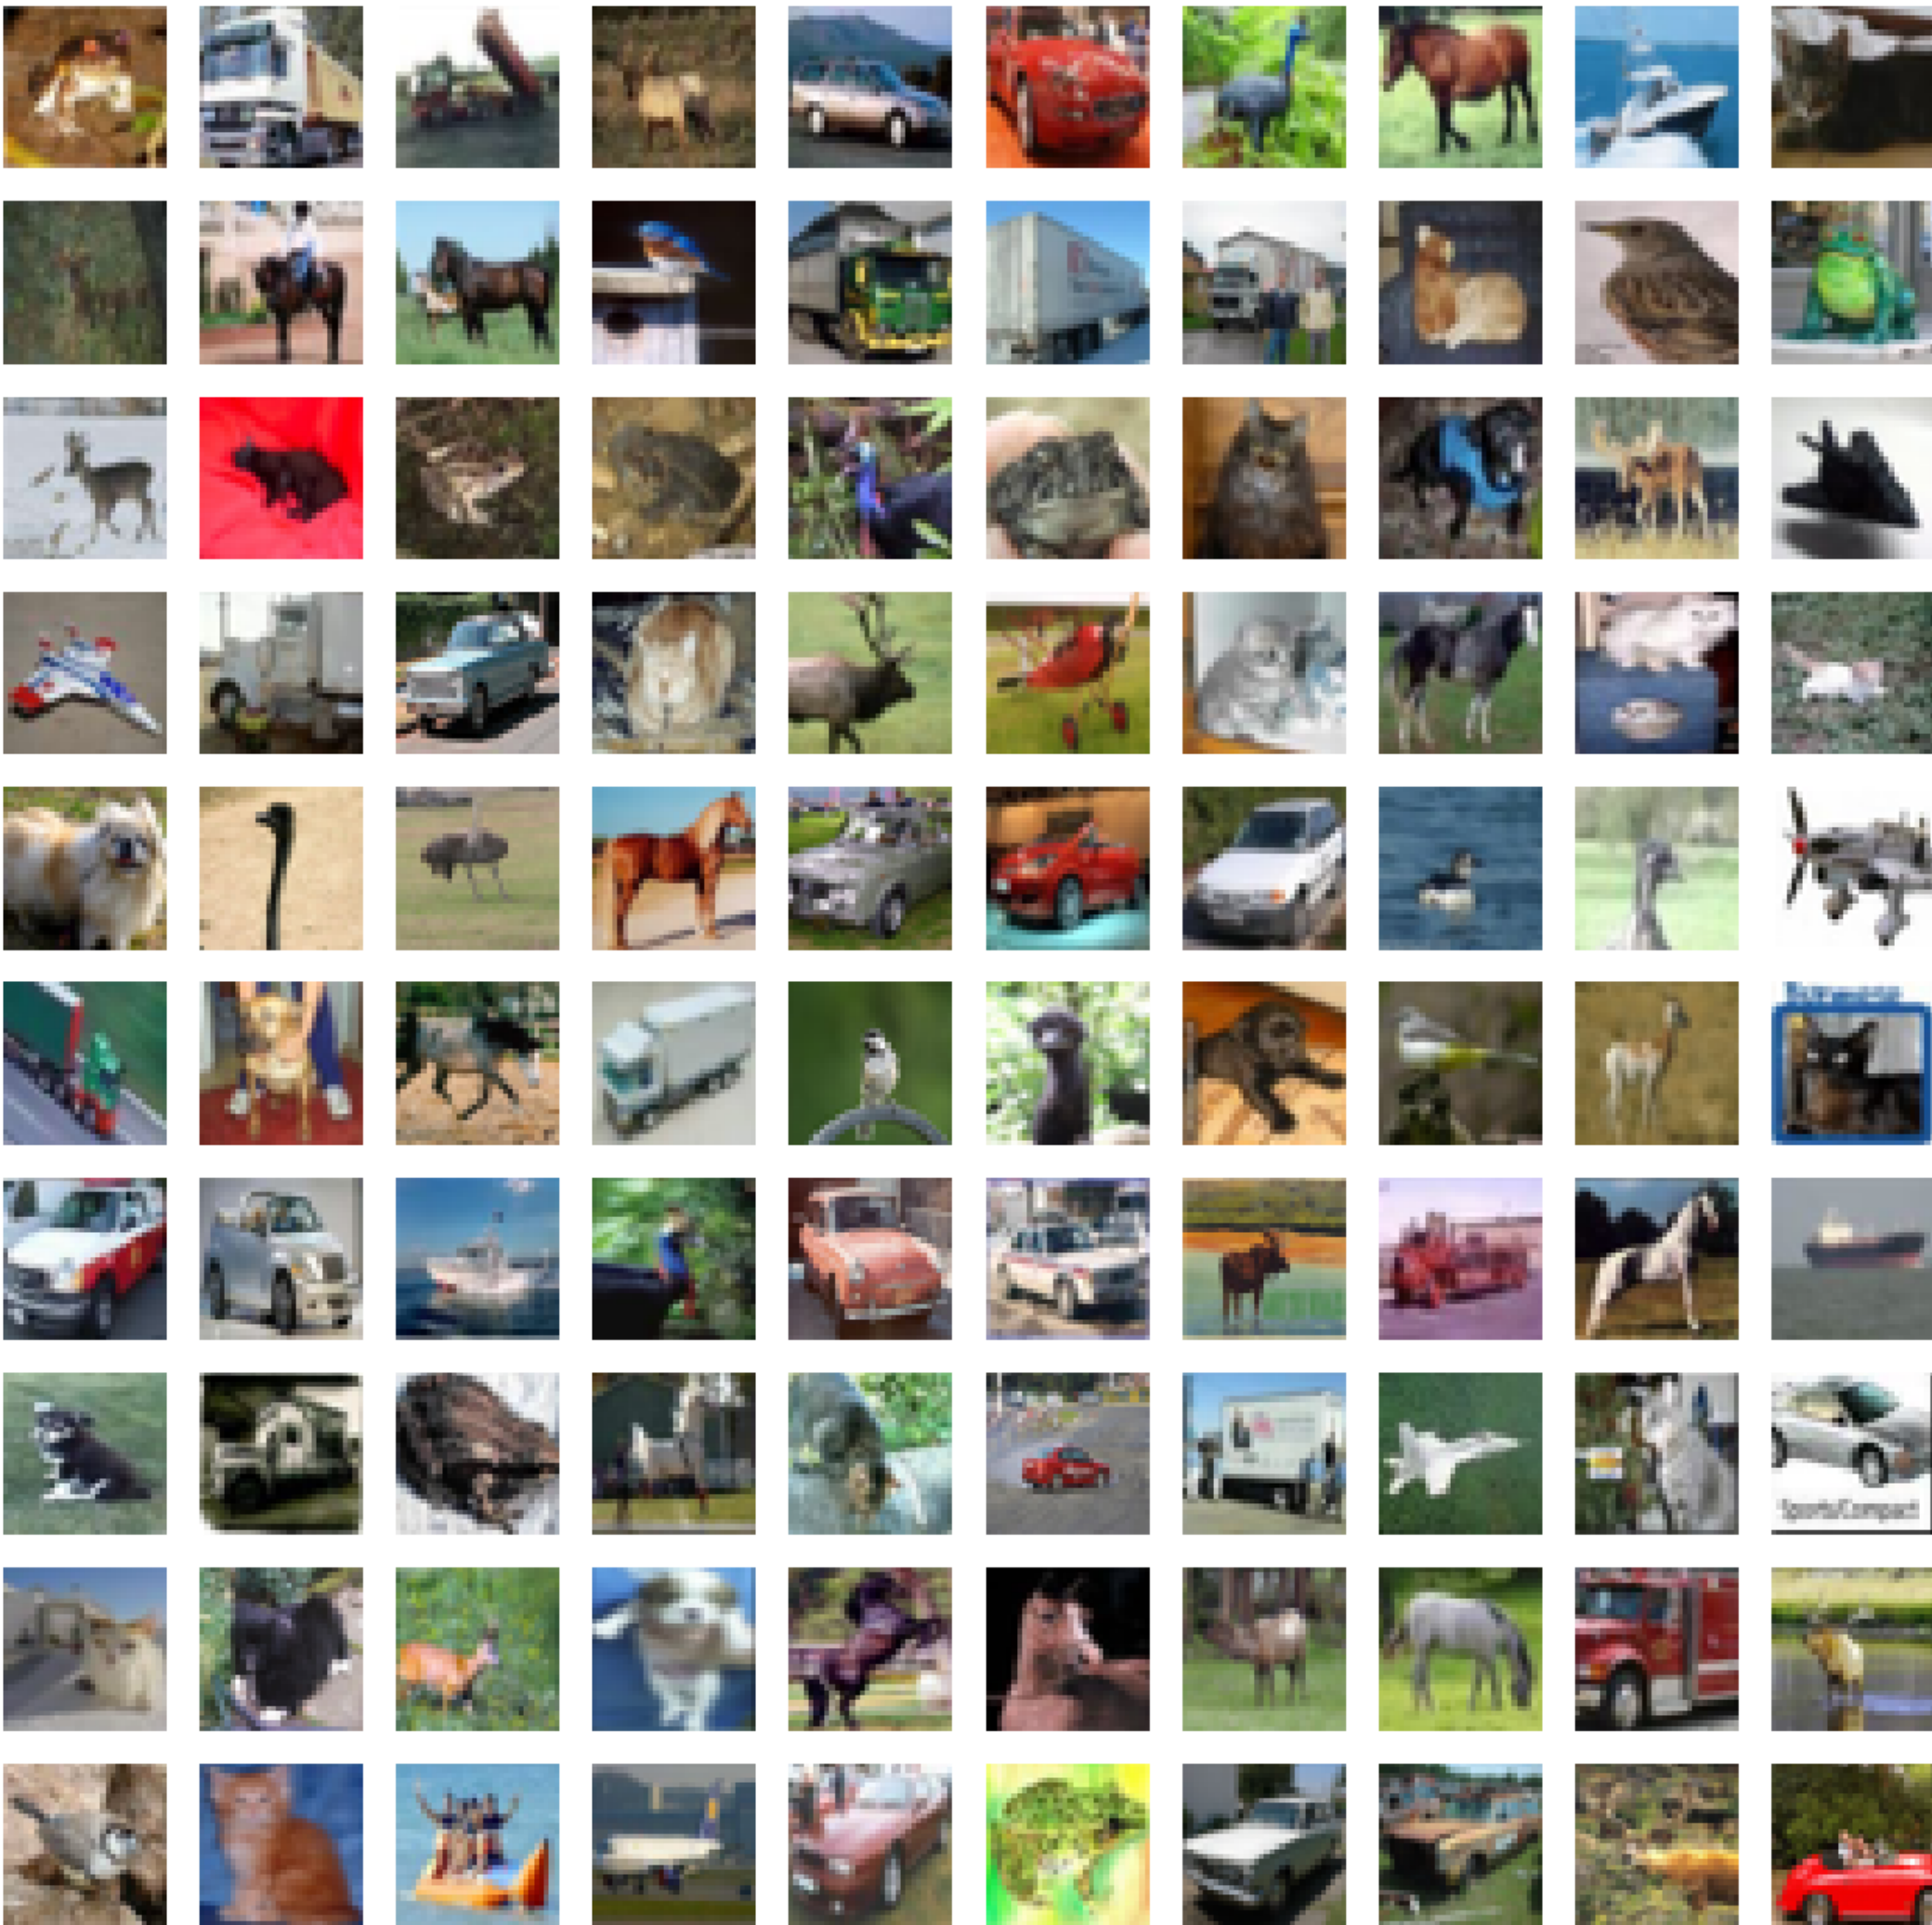
1 plt.figure(figsize = (16,16))
2 for i in range(100):
3     plt.subplot(10,10,1+i)
4     plt.imshow(train_images[i])

```

```

4 plt.axis('off')
5 plt.imshow(train_images[i], cmap = 'gray')

```



```

1 # Normalize pixel values to be between 0 and 1
2 train_images, test_images = train_images / 255.0, test_images / 255.0

```

```

1 from tensorflow.keras.models import Sequential
2 from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
3 from tensorflow.keras.activations import relu, sigmoid, tanh
4 from tensorflow.keras.optimizers import Adam
5 from tensorflow.keras import datasets, layers, models

```

```

1 # Define the CNN model
2 model = models.Sequential()
3 model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
4 model.add(layers.Conv2D(64, (3, 3), activation='relu'))
5 model.add(layers.MaxPooling2D((2, 2)))

```

```

6
7 # Add fully connected layers
8 model.add(layers.Flatten())
9 model.add(layers.Dense(64, activation='relu'))
10 model.add(layers.Dense(10))

```

```

1 # Compile the model
2 model.compile(optimizer='adam',
3               loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
4               metrics=['accuracy'])

```

```

1 # Train the model
2 history = model.fit(train_images, train_labels, epochs=5,
3                     validation_data=(test_images, test_labels))

```

```

Epoch 1/5
1563/1563 [=====] - 19s 5ms/step - loss: 1.3589 - accuracy: 0.5176 - val_loss: 1.1224 - val_accuracy: 0.6098
Epoch 2/5
1563/1563 [=====] - 9s 6ms/step - loss: 0.9895 - accuracy: 0.6563 - val_loss: 0.9849 - val_accuracy: 0.6615
Epoch 3/5
1563/1563 [=====] - 8s 5ms/step - loss: 0.8593 - accuracy: 0.6997 - val_loss: 0.9683 - val_accuracy: 0.6680
Epoch 4/5
1563/1563 [=====] - 8s 5ms/step - loss: 0.7516 - accuracy: 0.7395 - val_loss: 0.9347 - val_accuracy: 0.6841
Epoch 5/5
1563/1563 [=====] - 8s 5ms/step - loss: 0.6571 - accuracy: 0.7717 - val_loss: 0.9518 - val_accuracy: 0.6810

```

```

1 # Evaluate the model
2 test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
3 print(f"Test accuracy: {test_acc}")


313/313 - 1s - loss: 0.9518 - accuracy: 0.6810 - 702ms/epoch - 2ms/step
Test accuracy: 0.6809999942779541

```

```

1 # Plot training history
2 plt.plot(history.history['accuracy'], label='accuracy')
3 plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
4 plt.xlabel('Epoch')
5 plt.ylabel('Accuracy')
6 plt.ylim([0, 1])
7 plt.legend(loc='lower right')
8 plt.show()

```



**After simple CNN, we got 76% training accuracy and 69% testing accuracy**

▼ **2) Apply different number of convolutional layer + kernel size + channel number to progress good accuracy.**

**Adding 3 Conv Layer**

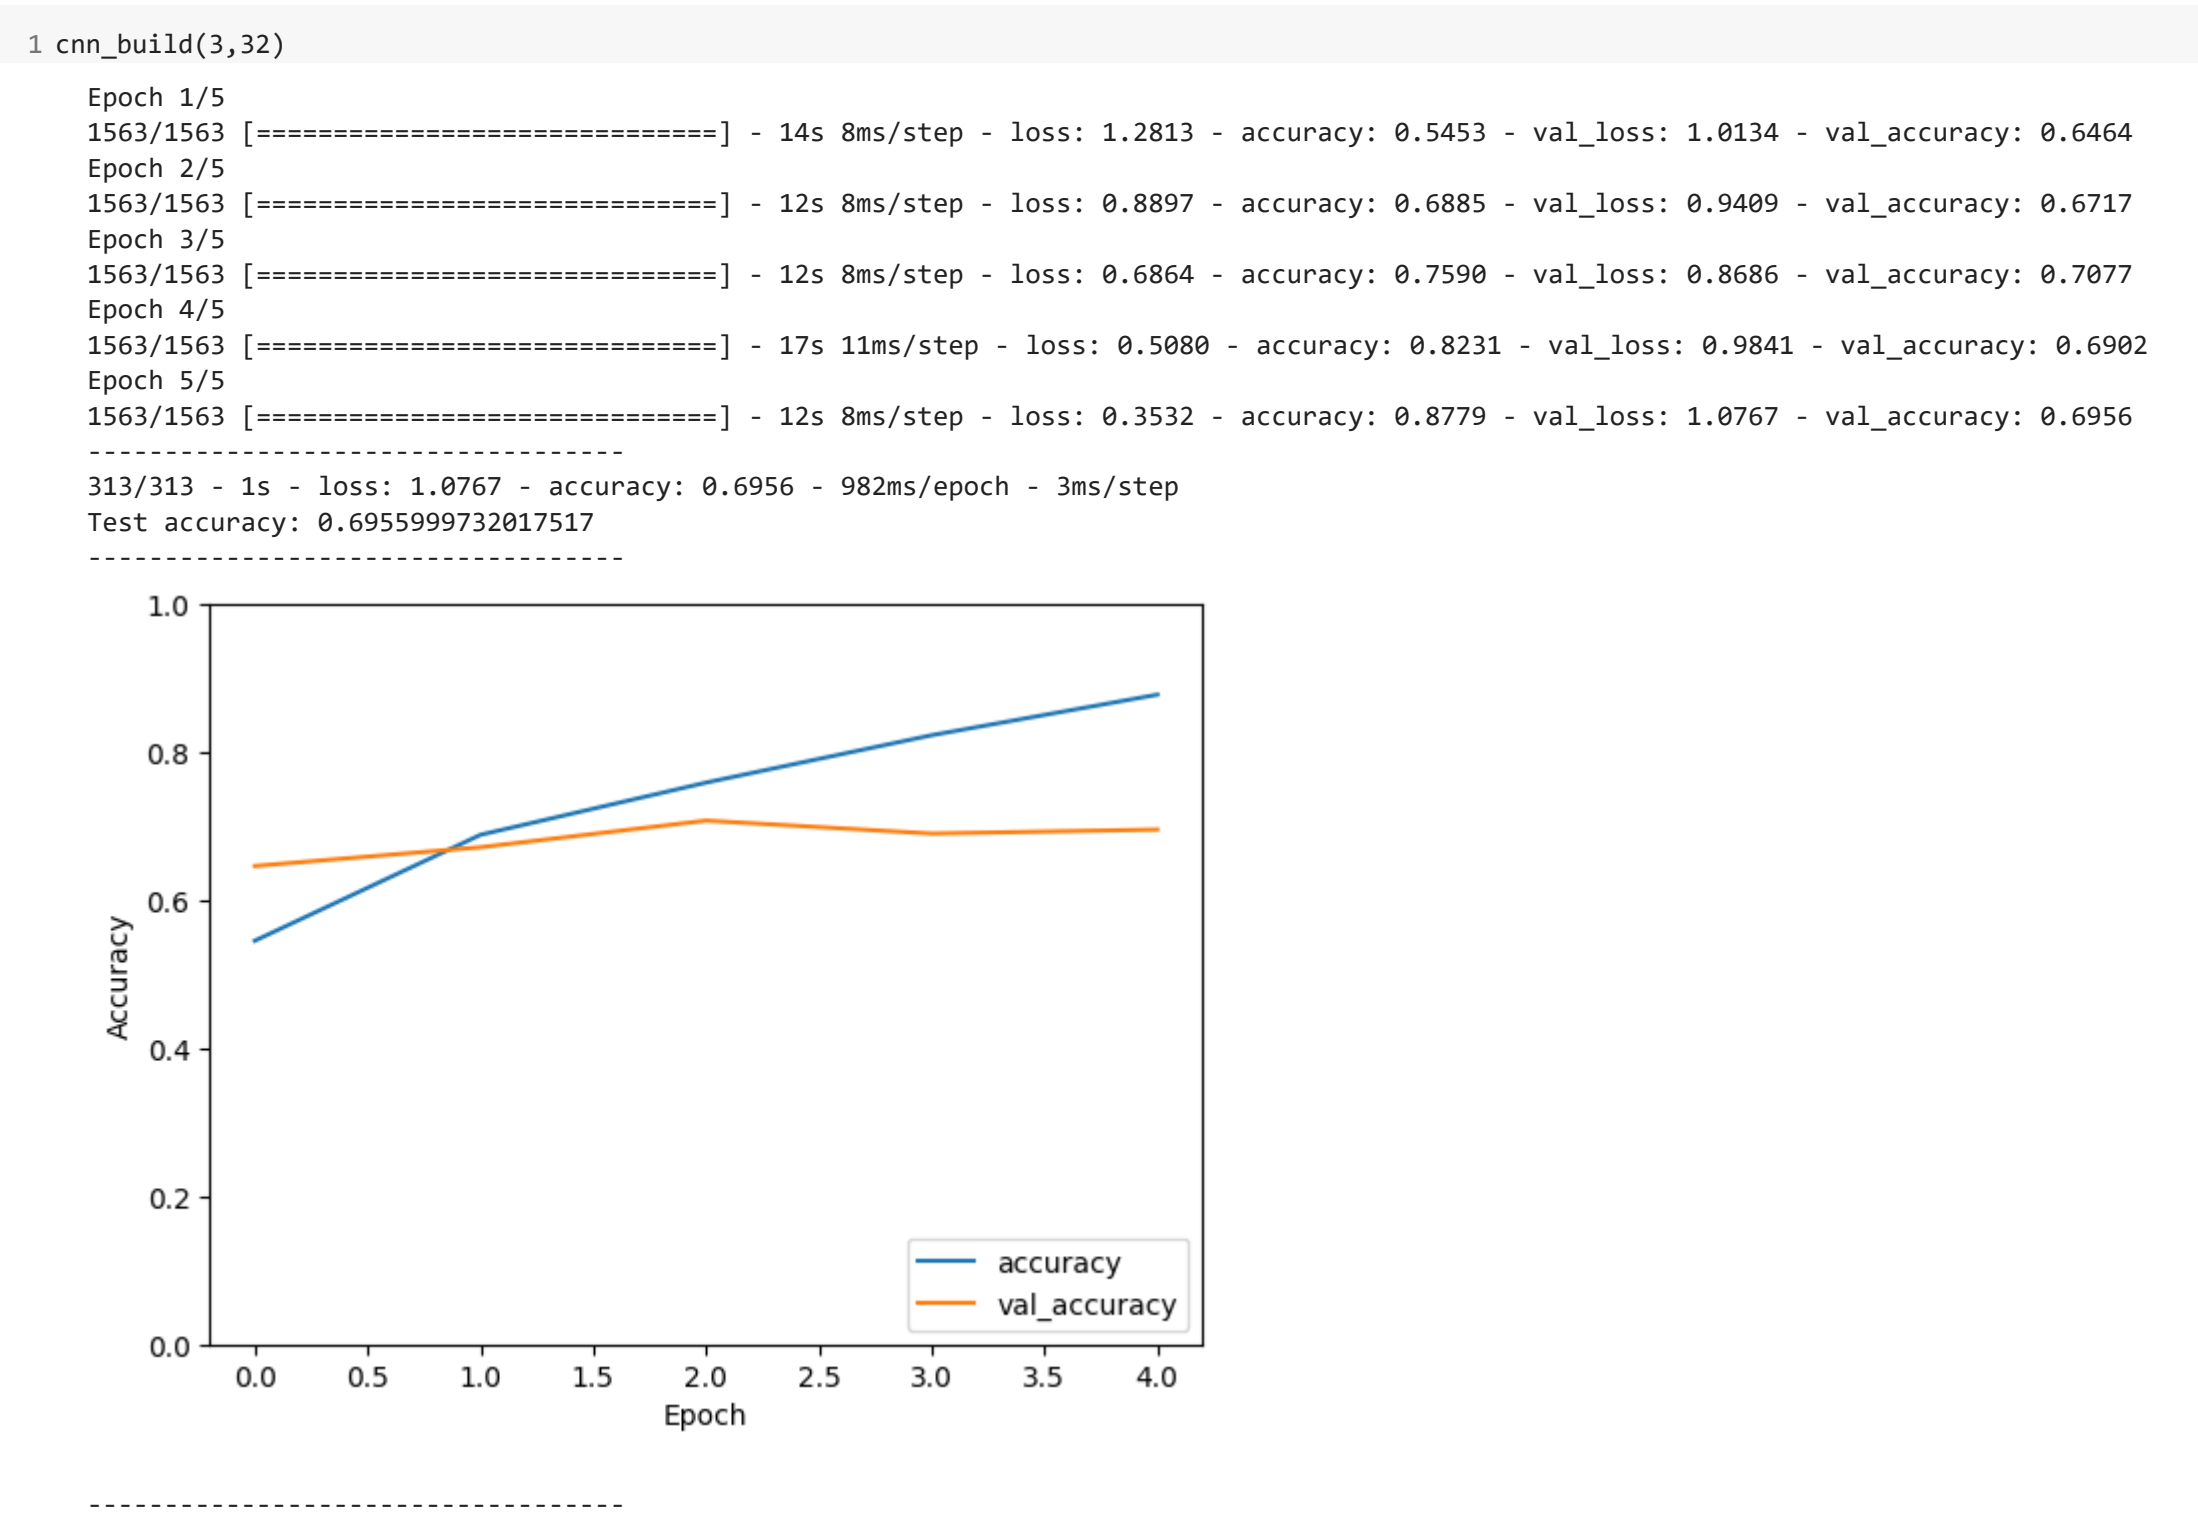
```

1 def cnn_build(kernel_size, channel):
2     # Define the CNN model
3     model = models.Sequential()
4
5     # Convolutional layers with varying parameters
6     model.add(layers.Conv2D(channel, (kernel_size, kernel_size), activation='relu', input_shape=(32, 32, 3)))
7     model.add(layers.Conv2D(channel*2, (kernel_size, kernel_size), activation='relu'))
8     model.add(layers.Conv2D(channel*4, (kernel_size, kernel_size), activation='relu'))

```

```
9 model.add(layers.MaxPooling2D((2, 2)))
10
11 # Flatten and add fully connected layers
12 model.add(layers.Flatten())
13 model.add(layers.Dense(128, activation='relu'))
14 model.add(layers.Dense(10))
15
16 # Compile the model
17 model.compile(optimizer='adam',
18               loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
19               metrics=['accuracy'])
20
21 # Train the model
22 history = model.fit(train_images, train_labels, epochs=5,
23                     validation_data=(test_images, test_labels))
24 print("-----")
25 # Evaluate the model
26 test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
27 print(f"Test accuracy: {test_acc}")
28 print("-----")
29
30 # Plot training history
31 plt.plot(history.history['accuracy'], label='accuracy')
32 plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
33 plt.xlabel('Epoch')
34 plt.ylabel('Accuracy')
35 plt.ylim([0, 1])
36 plt.legend(loc='lower right')
37 plt.show()
38 print()
39 print("-----")
40 print()
```

### Using 3 kernel size and 32 channels




### Using 5 kernel size and 48 channels

```
1 cnn_build(5,48)
```

Epoch 1/5  
1563/1563 [=====] - 18s 10ms/step - loss: 1.4937 - accuracy: 0.4570 - val loss: 1.2547 - val accuracy: 0.5581


Epoch 2/5  
1563/1563 [=====] - 15s 10ms/step - loss: 1.0963 - accuracy: 0.6138 - val\_loss: 1.0597 - val\_accuracy: 0.6320  
Epoch 3/5  
1563/1563 [=====] - 16s 10ms/step - loss: 0.8837 - accuracy: 0.6905 - val\_loss: 1.0614 - val\_accuracy: 0.6310  
Epoch 4/5  
1563/1563 [=====] - 15s 10ms/step - loss: 0.7137 - accuracy: 0.7501 - val\_loss: 1.0310 - val\_accuracy: 0.6680  
Epoch 5/5  
1563/1563 [=====] - 15s 10ms/step - loss: 0.5639 - accuracy: 0.8019 - val\_loss: 1.0581 - val\_accuracy: 0.6710  
-----  
313/313 - 1s - loss: 1.0581 - accuracy: 0.6710 - 1s/epoch - 3ms/step  
Test accuracy: 0.671000038146973  
-----



**Using 5 kernel size and 32 channels**

```
1 cnn_build(5,32)
```

Epoch 1/5  
1563/1563 [=====] - 15s 8ms/step - loss: 1.4813 - accuracy: 0.4638 - val\_loss: 1.2233 - val\_accuracy: 0.5633  
Epoch 2/5  
1563/1563 [=====] - 11s 7ms/step - loss: 1.0850 - accuracy: 0.6170 - val\_loss: 1.0533 - val\_accuracy: 0.6326  
Epoch 3/5  
1563/1563 [=====] - 11s 7ms/step - loss: 0.8785 - accuracy: 0.6920 - val\_loss: 1.0097 - val\_accuracy: 0.6526  
Epoch 4/5  
1563/1563 [=====] - 11s 7ms/step - loss: 0.7134 - accuracy: 0.7483 - val\_loss: 1.0447 - val\_accuracy: 0.6537  
Epoch 5/5  
1563/1563 [=====] - 11s 7ms/step - loss: 0.5507 - accuracy: 0.8083 - val\_loss: 1.1146 - val\_accuracy: 0.6575  
-----  
313/313 - 1s - loss: 1.1146 - accuracy: 0.6575 - 1s/epoch - 4ms/step  
Test accuracy: 0.6575000286102295  
-----




### Using 3 kernel size and 48 channels

```
1 cnn_build(3,48)
```

```
Epoch 1/5
1563/1563 [=====] - 19s 11ms/step - loss: 1.3856 - accuracy: 0.5018 - val_loss: 1.0578 - val_accuracy: 0.6252
Epoch 2/5
1563/1563 [=====] - 17s 11ms/step - loss: 0.9465 - accuracy: 0.6669 - val_loss: 0.9083 - val_accuracy: 0.6858
Epoch 3/5
1563/1563 [=====] - 17s 11ms/step - loss: 0.7563 - accuracy: 0.7362 - val_loss: 0.8792 - val_accuracy: 0.6948
Epoch 4/5
1563/1563 [=====] - 16s 11ms/step - loss: 0.5969 - accuracy: 0.7902 - val_loss: 0.8915 - val_accuracy: 0.7073
Epoch 5/5
1563/1563 [=====] - 17s 11ms/step - loss: 0.4403 - accuracy: 0.8458 - val_loss: 0.9679 - val_accuracy: 0.7058
-----
```

```
313/313 - 1s - loss: 0.9679 - accuracy: 0.7058 - 1s/epoch - 4ms/step
Test accuracy: 0.7057999968528748
-----
```



**Among all 3 Conv Layer with (33) kernel size and 48 channel performed best\***

**3) Implement the CNN model with different type of pooling method and discuss the interpretation happen due to these changes.**

### Using AveragePooling

```
1 # Define the CNN model
2 model = models.Sequential()
3
4 # Convolutional layers with varying parameters
5 model.add(layers.Conv2D(48, (3, 3), activation='relu', input_shape=(32, 32, 3)))
6 model.add(layers.Conv2D(48*2, (3, 3), activation='relu'))
7 model.add(layers.Conv2D(48*4, (3, 3), activation='relu'))
8 model.add(layers.AveragePooling2D((2, 2)))
9
10 # Flatten and add fully connected layers
11 model.add(layers.Flatten())
12 model.add(layers.Dense(128, activation='relu'))
13 model.add(layers.Dense(10))
14
15 # Compile the model
16 model.compile(optimizer='adam',
17               loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
18               metrics=['accuracy'])
19
20 # Train the model
21 history = model.fit(train_images, train_labels, epochs=5,
22                   validation_data=(test_images, test_labels))
23 print("-----")
24 # Evaluate the model
25 test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
26 print(f"Test accuracy: {test_acc}")
27 print("-----")
28
```



```


29 # Plot training history
30 plt.plot(history.history['accuracy'], label='accuracy')
31 plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
32 plt.xlabel('Epoch')
33 plt.ylabel('Accuracy')
34 plt.ylim([0, 1])
35 plt.legend(loc='lower right')
36 plt.show()
37 print()
38 print("-----")
39 print()

```

```

Epoch 1/5
1563/1563 [=====] - 19s 11ms/step - loss: 1.3088 - accuracy: 0.5317 - val_loss: 1.0279 - val_accuracy: 0.6413
Epoch 2/5
1563/1563 [=====] - 17s 11ms/step - loss: 0.8853 - accuracy: 0.6890 - val_loss: 0.9018 - val_accuracy: 0.6929
Epoch 3/5
1563/1563 [=====] - 17s 11ms/step - loss: 0.6866 - accuracy: 0.7609 - val_loss: 0.8482 - val_accuracy: 0.7090
Epoch 4/5
1563/1563 [=====] - 16s 11ms/step - loss: 0.5059 - accuracy: 0.8245 - val_loss: 0.8610 - val_accuracy: 0.7209
Epoch 5/5
1563/1563 [=====] - 17s 11ms/step - loss: 0.3485 - accuracy: 0.8777 - val_loss: 1.0453 - val_accuracy: 0.7166
-----
313/313 - 1s - loss: 1.0453 - accuracy: 0.7166 - 1s/epoch - 3ms/step
Test accuracy: 0.7166000008583069
-----

```



### Using GlobalAveragePooling

```

1 # Define the CNN model
2 model = models.Sequential()
3
4 # Convolutional layers with varying parameters
5 model.add(layers.Conv2D(48, (3, 3), activation='relu', input_shape=(32, 32, 3)))
6 model.add(layers.Conv2D(48*2, (3, 3), activation='relu'))
7 model.add(layers.Conv2D(48*4, (3, 3), activation='relu'))
8 model.add(layers.GlobalAveragePooling2D())
9
10 # Flatten and add fully connected layers
11 model.add(layers.Flatten())
12 model.add(layers.Dense(128, activation='relu'))
13 model.add(layers.Dense(10))
14
15 # Compile the model
16 model.compile(optimizer='adam',
17               loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
18               metrics=['accuracy'])
19
20 # Train the model
21 history = model.fit(train_images, train_labels, epochs=5,
22                    validation_data=(test_images, test_labels))
23 print("-----")
24 # Evaluate the model
25 test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
26 print(f"Test accuracy: {test_acc}")
27 print("-----")
28
29 # Plot training history

```

```


30 plt.plot(history.history['accuracy'], label='accuracy')
31 plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
32 plt.xlabel('Epoch')
33 plt.ylabel('Accuracy')
34 plt.ylim([0, 1])
35 plt.legend(loc='lower right')
36 plt.show()
37 print()
38 print("-----")
39 print()

```

```

Epoch 1/5
1563/1563 [=====] - 19s 11ms/step - loss: 1.7723 - accuracy: 0.3211 - val_loss: 1.5761 - val_accuracy: 0.3998
Epoch 2/5
1563/1563 [=====] - 16s 10ms/step - loss: 1.5014 - accuracy: 0.4427 - val_loss: 1.4114 - val_accuracy: 0.4767
Epoch 3/5
1563/1563 [=====] - 16s 10ms/step - loss: 1.3702 - accuracy: 0.5031 - val_loss: 1.3596 - val_accuracy: 0.5037
Epoch 4/5
1563/1563 [=====] - 16s 10ms/step - loss: 1.2836 - accuracy: 0.5385 - val_loss: 1.2076 - val_accuracy: 0.5669
Epoch 5/5
1563/1563 [=====] - 16s 10ms/step - loss: 1.2139 - accuracy: 0.5629 - val_loss: 1.1538 - val_accuracy: 0.5896
-----
313/313 - 1s - loss: 1.1538 - accuracy: 0.5896 - 1s/epoch - 3ms/step
Test accuracy: 0.5896000266075134
-----

```



***Among all the Pooling Layer, 'AveragePooling' performed best***

#### ***4) Apply different number of Stride and Padding techniques used for good result in CNN.***

##### ***Including Strides and Padding***

```

1 # Define the CNN model
2 model = models.Sequential()
3
4 # Convolutional layers with varying parameters
5 model.add(layers.Conv2D(48, (3, 3), activation='relu',padding='same', strides=(2, 2), input_shape=(32, 32, 3)))
6 model.add(layers.Conv2D(48*2, (3, 3),padding='same', strides=(2, 2), activation='relu'))
7 model.add(layers.Conv2D(48*4, (3, 3),padding='same', strides=(2, 2), activation='relu'))
8 model.add(layers.GlobalAveragePooling2D())
9
10 # Flatten and add fully connected layers
11 model.add(layers.Flatten())
12 model.add(layers.Dense(128, activation='relu'))
13 model.add(layers.Dense(10))
14
15 # Compile the model
16 model.compile(optimizer='adam',
17               loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
18               metrics=['accuracy'])
19
20 # Train the model
21 history = model.fit(train_images, train_labels, epochs=5,
22                    validation_data=(test_images, test_labels))
23 print("-----")

```



```


24 # Evaluate the model
25 test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
26 print(f"Test accuracy: {test_acc}")
27 print("-----")
28
29 # Plot training history
30 plt.plot(history.history['accuracy'], label='accuracy')
31 plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
32 plt.xlabel('Epoch')
33 plt.ylabel('Accuracy')
34 plt.ylim([0, 1])
35 plt.legend(loc='lower right')
36 plt.show()
37 print()
38 print("-----")
39 print()

```

```

Epoch 1/5
1563/1563 [=====] - 12s 6ms/step - loss: 1.7360 - accuracy: 0.3520 - val_loss: 1.5033 - val_accuracy: 0.4390
Epoch 2/5
1563/1563 [=====] - 9s 5ms/step - loss: 1.4207 - accuracy: 0.4789 - val_loss: 1.3356 - val_accuracy: 0.5089
Epoch 3/5
1563/1563 [=====] - 9s 5ms/step - loss: 1.2681 - accuracy: 0.5384 - val_loss: 1.2497 - val_accuracy: 0.5447
Epoch 4/5
1563/1563 [=====] - 11s 7ms/step - loss: 1.1666 - accuracy: 0.5774 - val_loss: 1.1676 - val_accuracy: 0.5836
Epoch 5/5
1563/1563 [=====] - 9s 6ms/step - loss: 1.0848 - accuracy: 0.6091 - val_loss: 1.0811 - val_accuracy: 0.6128
-----
313/313 - 1s - loss: 1.0811 - accuracy: 0.6128 - 719ms/epoch - 2ms/step
Test accuracy: 0.612800020980835
-----

```



```

1 # Define the CNN model
2 model = models.Sequential()
3
4 # Convolutional layers with varying parameters
5 model.add(layers.Conv2D(48, (3, 3), activation='relu',padding='same', strides=(4, 4), input_shape=(32, 32, 3)))
6 model.add(layers.Conv2D(48*2, (3, 3),padding='same', strides=(4, 4), activation='relu'))
7 model.add(layers.Conv2D(48*4, (3, 3),padding='same', strides=(4, 4), activation='relu'))
8 model.add(layers.GlobalAveragePooling2D())
9
10 # Flatten and add fully connected layers
11 model.add(layers.Flatten())
12 model.add(layers.Dense(128, activation='relu'))
13 model.add(layers.Dense(10))
14
15 # Compile the model
16 model.compile(optimizer='adam',
17               loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
18               metrics=['accuracy'])
19
20 # Train the model
21 history = model.fit(train_images, train_labels, epochs=5,
22                    validation_data=(test_images, test_labels))
23 print("-----")
24 # Evaluate the model
25 test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
26 print(f"Test accuracy: {test_acc}")
27 print("-----")

```

```


28
29 # Plot training history
30 plt.plot(history.history['accuracy'], label='accuracy')
31 plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
32 plt.xlabel('Epoch')
33 plt.ylabel('Accuracy')
34 plt.ylim([0, 1])
35 plt.legend(loc='lower right')
36 plt.show()
37 print()
38 print("-----")
39 print()

```

```

Epoch 1/5
1563/1563 [=====] - 11s 6ms/step - loss: 1.7024 - accuracy: 0.3802 - val_loss: 1.5150 - val_accuracy: 0.4479
Epoch 2/5
1563/1563 [=====] - 8s 5ms/step - loss: 1.4434 - accuracy: 0.4776 - val_loss: 1.3705 - val_accuracy: 0.5066
Epoch 3/5
1563/1563 [=====] - 9s 6ms/step - loss: 1.3246 - accuracy: 0.5223 - val_loss: 1.2988 - val_accuracy: 0.5282
Epoch 4/5
1563/1563 [=====] - 9s 6ms/step - loss: 1.2357 - accuracy: 0.5542 - val_loss: 1.2773 - val_accuracy: 0.5380
Epoch 5/5
1563/1563 [=====] - 8s 5ms/step - loss: 1.1653 - accuracy: 0.5815 - val_loss: 1.2489 - val_accuracy: 0.5478
-----
313/313 - 1s - loss: 1.2489 - accuracy: 0.5478 - 946ms/epoch - 3ms/step
Test accuracy: 0.5478000044822693
-----

```



***Padding and incresing Stride decreses Accuracy, thus model is not able to capture patterns***

## 5) Adjust the CNN model with the various hyper parameter tuning (Optimizer, Activation function, epoch, Learning rate etc.)

***Instead of Adam using Adagrad Optimisers with 3 epoch***

```

1 from tensorflow.keras.optimizers import Adagrad

```

```

1 def build_cnn(activation):
2     model = Sequential([
3         Conv2D(32, (3, 3), activation=activation, input_shape=(32, 32, 3)),
4         Conv2D(64, (3, 3), activation=activation),
5         MaxPooling2D((2, 2)),
6         Flatten(),
7         Dense(64, activation=activation),
8         Dense(10, activation='softmax')
9     ])
10    return model
11 data = []

```

```

1 from sklearn.metrics import confusion_matrix, recall_score, f1_score, precision_score, classification_report
2
3 def train_and_evaluate(learning_rate, activation):
4     model = build_cnn(activation)
5     model.compile(optimizer=Adagrad(learning_rate=learning_rate),

```

```

6         loss='sparse_categorical_crossentropy',
7         metrics=['accuracy'])
8
9     history = model.fit(train_images, train_labels, epochs=3, validation_split=0.2)
10    l =[]
11    # Evaluate the model
12    test_loss, test_acc = model.evaluate(test_images, test_labels)
13    y_pred = np.argmax(model.predict(test_images), axis=-1)
14
15    print("Test accuracy:", test_acc)
16    print("Confusion matrix:\n", confusion_matrix(test_labels, y_pred))
17    print("Classification report:\n", classification_report(test_labels, y_pred))
18    precision = precision_score(test_labels, y_pred, average='weighted')
19    recall = recall_score(test_labels, y_pred, average='weighted')
20    f1 = f1_score(test_labels, y_pred, average='weighted')
21    l = [learning_rate,activation,test_acc,confusion_matrix(test_labels, y_pred),classification_report(test_labels, y_pred),precision,recall)
22    data.append(l)
23    print("Precision:", precision)
24    print("Recall:", recall)
25    print("F1-score:", f1)
26    # Plot learning curves
27    plt.plot(history.history['accuracy'], label='accuracy')
28    plt.plot(history.history['val_accuracy'], label='val_accuracy')
29    plt.xlabel('Epoch')
30    plt.ylabel('Accuracy')
31    plt.legend()
32    plt.show()

```

```

1 learning_rates = [0.001, 0.01, 0.1]
2 activation_functions = [relu, sigmoid, tanh]

```

```

1 for lr in learning_rates:
2     for activation in activation_functions:
3         print(f"Learning rate: {lr}, Activation: {activation.__name__}")
4         train_and_evaluate(lr, activation)
5         print("-----")

```

```

Learning rate: 0.001, Activation: relu
Epoch 1/3
1250/1250 [=====] - 8s 5ms/step - loss: 2.1108 - accuracy: 0.2448 - val_loss: 1.9362 - val_accuracy: 0.3263
Epoch 2/3
1250/1250 [=====] - 5s 4ms/step - loss: 1.8493 - accuracy: 0.3584 - val_loss: 1.7954 - val_accuracy: 0.3762
Epoch 3/3
1250/1250 [=====] - 7s 6ms/step - loss: 1.7454 - accuracy: 0.3943 - val_loss: 1.7215 - val_accuracy: 0.4005
313/313 [=====] - 1s 3ms/step - loss: 1.7091 - accuracy: 0.4116
313/313 [=====] - 1s 2ms/step
Test accuracy: 0.411599937057495
Confusion matrix:
[[486  77  29  45  27  31  24  34 211  36]
 [ 49 641   2  35   5  31  58  29  62  88]
 [112  58 120 117 242 131  96  62  45  17]
 [ 31  93  32 332  59 217 111  59  33  33]
 [ 51  38  43 110 383 106 138  90  31  10]
 [ 19  53  47 200  82 398  70  79  38  14]
 [  5  56  33 186 131  80 446  21  14  28]
 [ 40  68  13 141  85 116  41 409  34  53]
 [139 121   7  36   4  44   7  22 555  65]
 [ 55 322   4  46   7  14  56  41 109 346]]
Classification report:


```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.49      | 0.49   | 0.49     | 1000    |
| 1            | 0.42      | 0.64   | 0.51     | 1000    |
| 2            | 0.36      | 0.12   | 0.18     | 1000    |
| 3            | 0.27      | 0.33   | 0.30     | 1000    |
| 4            | 0.37      | 0.38   | 0.38     | 1000    |
| 5            | 0.34      | 0.40   | 0.37     | 1000    |
| 6            | 0.43      | 0.45   | 0.44     | 1000    |
| 7            | 0.48      | 0.41   | 0.44     | 1000    |
| 8            | 0.49      | 0.56   | 0.52     | 1000    |
| 9            | 0.50      | 0.35   | 0.41     | 1000    |
| accuracy     |           |        | 0.41     | 10000   |
| macro avg    | 0.42      | 0.41   | 0.40     | 10000   |
| weighted avg | 0.42      | 0.41   | 0.40     | 10000   |

```


Precision: 0.415741500800354
Recall: 0.4116
F1-score: 0.4026740934175066

```





***Analysing all the results, we can conclude that ADAM optimiser with RELU as activation function having 33 kernel size and 48 channels performed best among all the combinations.\****

[illegible]