

Deep Learning: Assignment-8

Train a GAN Model on any dataset

Submitted By: Mrinal Bhan

DSAI 211020428

Importing the libraries

```
1 from __future__ import print_function
2 import time
3 import torch
4 import torch.nn as nn
5 import torch.nn.parallel
6 import torch.optim as optim
7 import torch.utils.data
8 import torchvision.datasets as dset
9 import torchvision.transforms as transforms
10 import torchvision.utils as vutils
11 from torch.autograd import Variable
12 import matplotlib.pyplot as plt
13 import numpy as np
14 from torch import nn, optim
15 import torch.nn.functional as F
16 from torchvision import datasets, transforms
17 from torchvision.utils import save_image
18 import matplotlib.pyplot as plt
19 import matplotlib.image as mpimg
20 from tqdm import tqdm_notebook as tqdm
```

For this assignment, We've used The Stanford Dogs dataset Context

The Stanford Dogs dataset contains images of 120 breeds of dogs from around the world. This dataset has been built using images and annotation from ImageNet for the task of fine-grained image categorization. It was originally collected for fine-grain image categorization, a challenging problem as certain dog breeds have near identical features or differ in colour and age.

Content:

Number of categories: 120

Number of images: 20,580

Annotations: Class labels, Bounding boxes


```
1 PATH = '../input/all-dogs/all-dogs/'
2 images = os.listdir(PATH)
3 print(f'There are {len(os.listdir(PATH))} pictures of dogs.')
4
5 fig, axes = plt.subplots(nrows=3, ncols=3, figsize=(12,10))
6
7 for indx, axis in enumerate(axes.flatten()):
8     rnd_indx = np.random.randint(0, len(os.listdir(PATH)))
9     img = plt.imread(PATH + images[rnd_indx])
10    imgplot = axis.imshow(img)
11    axis.set_title(images[rnd_indx])
12    axis.set_axis_off()
13 plt.tight_layout(rect=[0, 0.03, 1, 0.95])
```

There are 20579 pictures of dogs.

n02108089_5599.jpg




n02090379_2720.jpg



n02088632_205.jpg



n02096051_6972.jpg



n02093859_508.jpg



n02105505_358.jpg





n02100583_1560.jpg



n02100583_6040.jpg



n02106166_1205.jpg



Pre-processing the images in the dataset

```
1 batch_size = 32
2 image_size = 64
3
4 random_transforms = [transforms.ColorJitter(), transforms.RandomRotation(degrees=20)]
5 transform = transforms.Compose([transforms.Resize(64),
6                                 transforms.CenterCrop(64),
7                                 transforms.RandomHorizontalFlip(p=0.5),
8                                 transforms.RandomApply(random_transforms, p=0.2),
9                                 transforms.ToTensor(),
10                                transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
11
12 train_data = datasets.ImageFolder('../input/all-dogs/', transform=transform)
13 train_loader = torch.utils.data.DataLoader(train_data, shuffle=True,
14                                             batch_size=batch_size)
15
16 imgs, label = next(iter(train_loader))
17 imgs = imgs.numpy().transpose(0, 2, 3, 1)
```

Initial Results

Defining the Weights

```
1 def weights_init(m):
2     """
3     Takes as input a neural network m that will initialize all its weights.
4     """
5     classname = m.__class__.__name__
6     if classname.find('Conv') != -1:
7         m.weight.data.normal_(0.0, 0.02)
8     elif classname.find('BatchNorm') != -1:
9         m.weight.data.normal_(1.0, 0.02)
10        m.bias.data.fill_(0)
```

Generator

```
1 class G(nn.Module):
2     def __init__(self):
3         # Used to inherit the torch.nn Module
4         super(G, self).__init__()
5         # Meta Module - consists of different layers of Modules
6         self.main = nn.Sequential(
7             nn.ConvTranspose2d(100, 512, 4, stride=1, padding=0, bias=False),
8             nn.BatchNorm2d(512),
9             nn.ReLU(True),
10            nn.ConvTranspose2d(512, 256, 4, stride=2, padding=1, bias=False),
11            nn.BatchNorm2d(256),
12            nn.ReLU(True),
13            nn.ConvTranspose2d(256, 128, 4, stride=2, padding=1, bias=False),
14            nn.BatchNorm2d(128),
15            nn.ReLU(True),
16            nn.ConvTranspose2d(128, 64, 4, stride=2, padding=1, bias=False),
17            nn.BatchNorm2d(64),
18            nn.ReLU(True),
19            nn.ConvTranspose2d(64, 3, 4, stride=2, padding=1, bias=False),
20            nn.Tanh()
21        )
22
23    def forward(self, input):
24        output = self.main(input)
25        return output
26
```

```

27 # Creating the generator
28 netG = G()
29 netG.apply(weights_init)

G(
  (main): Sequential(
    (0): ConvTranspose2d(100, 512, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU(inplace)
    (3): ConvTranspose2d(512, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (4): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (5): ReLU(inplace)
    (6): ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (7): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (8): ReLU(inplace)
    (9): ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (10): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (11): ReLU(inplace)
    (12): ConvTranspose2d(64, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (13): Tanh()
  )
)

```

Discriminator

```

1 # Defining the discriminator
2 class D(nn.Module):
3     def __init__(self):
4         super(D, self).__init__()
5         self.main = nn.Sequential(
6             nn.Conv2d(3, 64, 4, stride=2, padding=1, bias=False),
7             nn.LeakyReLU(negative_slope=0.2, inplace=True),
8             nn.Conv2d(64, 128, 4, stride=2, padding=1, bias=False),
9             nn.BatchNorm2d(128),
10            nn.LeakyReLU(negative_slope=0.2, inplace=True),
11            nn.Conv2d(128, 256, 4, stride=2, padding=1, bias=False),
12            nn.BatchNorm2d(256),
13            nn.LeakyReLU(negative_slope=0.2, inplace=True),
14            nn.Conv2d(256, 512, 4, stride=2, padding=1, bias=False),
15            nn.BatchNorm2d(512),
16            nn.LeakyReLU(negative_slope=0.2, inplace=True),
17            nn.Conv2d(512, 1, 4, stride=1, padding=0, bias=False),
18            nn.Sigmoid()
19        )
20
21    def forward(self, input):
22        output = self.main(input)
23        # .view(-1) = Flattens the output into 1D instead of 2D
24        return output.view(-1)
25
26
27 # Creating the discriminator
28 netD = D()
29 netD.apply(weights_init)
30

```

```

D(
  (main): Sequential(
    (0): Conv2d(3, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): LeakyReLU(negative_slope=0.2, inplace)
    (2): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (4): LeakyReLU(negative_slope=0.2, inplace)
    (5): Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (6): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (7): LeakyReLU(negative_slope=0.2, inplace)
    (8): Conv2d(256, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (9): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (10): LeakyReLU(negative_slope=0.2, inplace)
    (11): Conv2d(512, 1, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (12): Sigmoid()
  )
)

```

Training Parameters

```

1 batch_size = 32
2 LR_G = 0.001
3 LR_D = 0.0005
4
5 beta1 = 0.5
6 epochs = 100
7
8 real_label = 0.9

```

```

9 fake_label = 0
10 nz = 128
11
12 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

```

Initialize models and optimizers

```

1 netG = Generator(nz).to(device)
2 netD = Discriminator().to(device)
3
4 criterion = nn.BCELoss()
5
6 optimizerD = optim.Adam(netD.parameters(), lr=LR_D, betas=(beta1, 0.999))
7 optimizerG = optim.Adam(netG.parameters(), lr=LR_G, betas=(beta1, 0.999))
8
9 fixed_noise = torch.randn(25, nz, 1, 1, device=device)
10
11 G_losses = []
12 D_losses = []
13 epoch_time = []

```

```

1 def plot_loss (G_losses, D_losses, epoch):
2     plt.figure(figsize=(10,5))
3     plt.title("Generator and Discriminator Loss - EPOCH "+ str(epoch))
4     plt.plot(G_losses,label="G")
5     plt.plot(D_losses,label="D")
6     plt.xlabel("iterations")
7     plt.ylabel("Loss")
8     plt.legend()
9     plt.show()

```

Show generated images

```

1 def show_generated_img(n_images=5):
2     sample = []
3     for _ in range(n_images):
4         noise = torch.randn(1, nz, 1, 1, device=device)
5         gen_image = netG(noise).to("cpu").clone().detach().squeeze(0)
6         gen_image = gen_image.numpy().transpose(1, 2, 0)
7         sample.append(gen_image)
8
9     figure, axes = plt.subplots(1, len(sample), figsize = (64,64))
10    for index, axis in enumerate(axes):
11        axis.axis('off')
12        image_array = sample[index]
13        axis.imshow(image_array)
14
15    plt.show()
16    plt.close()

```

Training Loop

```

1 for epoch in range(epochs):
2
3     start = time.time()
4     for ii, (real_images, train_labels) in tqdm(enumerate(train_loader), total=len(train_loader)):
5         #####
6         # (1) Update D network: maximize log(D(x)) + log(1 - D(G(z)))
7         #####
8         # train with real
9         netD.zero_grad()
10        real_images = real_images.to(device)
11        batch_size = real_images.size(0)
12        labels = torch.full((batch_size, 1), real_label, device=device)
13
14        output = netD(real_images)
15        errD_real = criterion(output, labels)
16        errD_real.backward()
17        D_x = output.mean().item()
18
19        # train with fake
20        noise = torch.randn(batch_size, nz, 1, 1, device=device)
21        fake = netG(noise)
22        labels.fill_(fake_label)
23        output = netD(fake.detach())
24        errD_fake = criterion(output, labels)
25        errD_fake.backward()
26        D_G_z1 = output.mean().item()
27        errD = errD_real + errD_fake
28        optimizerD.step()

```




```
29 #####
30 # (2) Update G network: maximize log(D(G(z)))
31 #####
32 netG.zero_grad()
33 labels.fill_(real_label) # fake labels are real for generator cost
34 output = netD(fake)
35 errG = criterion(output, labels)
36 errG.backward()
37 D_G_z2 = output.mean().item()
38 optimizerG.step()
39
40
41 # Save Losses for plotting later
42 G_losses.append(errG.item())
43 D_losses.append(errD.item())
44
45 if (ii+1) % (len(train_loader)//2) == 0:
46     print('%d/%d [%d/%d] Loss_D: %.4f Loss_G: %.4f D(x): %.4f D(G(z)): %.4f / %.4f'
47           % (epoch + 1, epochs, ii+1, len(train_loader),
48             errD.item(), errG.item(), D_x, D_G_z1, D_G_z2))
49
50 plot_loss (G_losses, D_losses, epoch)
51 G_losses = []
52 D_losses = []
53 if epoch % 10 == 0:
54     show_generated_img()
55
56 epoch_time.append(time.time()- start)
57
58 #         valid_image = netG(fixed_noise)
```

```
1 print (">> average EPOCH duration = ", np.mean(epoch_time))
>> average EPOCH duration = 121.24567284107208
```

Generated Images

```
1 show_generated_img(7)
```



```
1 fig = plt.figure(figsize=(25, 16))
2 # display 10 images from each class
3 for i, j in enumerate(images[:32]):
4     ax = fig.add_subplot(4, 8, i + 1, xticks=[], yticks=[])
5     plt.imshow(j)
```

