

Deep Learning: Assignment-9

Training a LSTM Model to predict stock prices of COAL India Stock price

Submitted By: Mrinal Bhan

DSAI 211020428

1.Download/search the Stock market dataset from <https://www.kaggle.com/datasets/rohanrao/nifty50-stock-market-data>. You have given the flexibility to use any other stock dataset if found suitable.

```
1 import numpy as np
2 import seaborn as sns
3 import pandas as pd
4 import matplotlib.pyplot as plt
5 from datetime import datetime as dt
6
```

```
1 data = pd.read_csv('/kaggle/input/nifty50-stock-market-data/COALINDIA.csv')
```

```
1 data.head(10)
```

	Date	Symbol	Series	Prev Close	Open	High	Low	Last	Close	VWAP	Volume	Turnover	Trades	Deliverable Volume	%Deliv
0	2010-11-04	COALINDIA	EQ	245.00	291.00	344.90	291.00	342.00	342.55	327.29	479716245	1.570040e+16	NaN	187584905	0
1	2010-11-05	COALINDIA	EQ	342.55	343.00	356.50	343.00	348.30	349.85	349.78	31927173	1.116747e+15	NaN	10894509	0
2	2010-11-08	COALINDIA	EQ	349.85	351.80	355.90	329.50	331.40	330.75	335.19	46932779	1.573118e+15	NaN	16651623	0
3	2010-11-09	COALINDIA	EQ	330.75	330.15	333.40	325.00	325.40	326.05	327.75	23741956	7.781383e+14	NaN	12977359	0
4	2010-11-10	COALINDIA	EQ	326.05	325.40	327.80	320.05	321.30	322.80	323.78	21057129	6.817982e+14	NaN	6280335	0
5	2010-11-11	COALINDIA	EQ	322.80	323.00	336.95	321.85	329.60	328.55	331.43	26548372	8.799018e+14	NaN	9365609	0
6	2010-11-12	COALINDIA	EQ	328.55	330.00	332.95	318.00	320.55	320.15	325.77	15004107	4.887872e+14	NaN	5797721	0
7	2010-11-15	COALINDIA	EQ	320.15	321.20	322.90	315.05	318.10	317.20	319.39	9917395	3.167544e+14	NaN	4812011	0
8	2010-11-16	COALINDIA	EQ	317.20	320.00	323.95	316.10	317.50	318.05	319.21	12100114	3.862460e+14	NaN	6620732	0
9	2010-11-18	COALINDIA	EQ	318.05	319.60	328.95	310.00	327.30	326.90	319.09	23983896	7.652941e+14	NaN	9152722	0

```
1 data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2598 entries, 0 to 2597
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  2598 non-null  object
1   Symbol                2598 non-null  object
2   Series                2598 non-null  object
3   Prev Close            2598 non-null  float64
4   Open                  2598 non-null  float64
5   High                  2598 non-null  float64
6   Low                   2598 non-null  float64
7   Last                  2598 non-null  float64
8   Close                 2598 non-null  float64
9   VWAP                  2598 non-null  float64
10  Volume                2598 non-null  int64
11  Turnover              2598 non-null  float64
12  Trades                2456 non-null  float64
13  Deliverable Volume    2598 non-null  int64
14  %Deliverble           2598 non-null  float64
dtypes: float64(10), int64(2), object(3)
memory usage: 304.6+ KB
```

2. Apply any two Preprocessing technique to the above datasets.

```
1 data.isnull().sum()
```

Date	0
Symbol	0
Series	0
Prev Close	0
Open	0
High	0
Low	0
Last	0
Close	0

VWAP 0
Volume 0
Turnover 0
Trades 142
Deliverable Volume 0
%Deliverble 0
dtype: int64

```
1 data = data.drop(['Symbol', 'Series', 'Prev Close', 'Last', 'VWAP', 'Turnover', 'Trades', 'Deliverable Volume', '%Deliverble'], axis = 1)
```

```
1 data.head(10)
```

	Date	Open	High	Low	Close	Volume
0	2010-11-04	291.00	344.90	291.00	342.55	479716245
1	2010-11-05	343.00	356.50	343.00	349.85	31927173
2	2010-11-08	351.80	355.90	329.50	330.75	46932779
3	2010-11-09	330.15	333.40	325.00	326.05	23741956
4	2010-11-10	325.40	327.80	320.05	322.80	21057129
5	2010-11-11	323.00	336.95	321.85	328.55	26548372
6	2010-11-12	330.00	332.95	318.00	320.15	15004107
7	2010-11-15	321.20	322.90	315.05	317.20	9917395
8	2010-11-16	320.00	323.95	316.10	318.05	12100114
9	2010-11-18	319.60	328.95	310.00	326.90	23983896


```
1 data.isnull().sum()
```

Date 0
Open 0
High 0
Low 0
Close 0
Volume 0
dtype: int64

```
1 plt.plot(data['Close'], label="Close price")  
2 plt.xlabel("Timestamp")  
3 plt.ylabel("Closing price")  
4 df = data  
5 print(df)
```

	Date	Open	High	Low	Close	Volume
0	2010-11-04	291.00	344.90	291.00	342.55	479716245
1	2010-11-05	343.00	356.50	343.00	349.85	31927173
2	2010-11-08	351.80	355.90	329.50	330.75	46932779
3	2010-11-09	330.15	333.40	325.00	326.05	23741956
4	2010-11-10	325.40	327.80	320.05	322.80	21057129
...
2593	2021-04-26	127.75	128.00	126.50	126.95	4967884
2594	2021-04-27	127.00	127.90	126.60	127.50	3982954
2595	2021-04-28	128.00	129.45	127.50	128.50	6206074
2596	2021-04-29	129.75	130.05	127.65	128.05	8345584
2597	2021-04-30	127.40	134.60	127.00	133.05	27396950

[2598 rows x 6 columns]



```
1 data.reset_index(drop=True, inplace=True)
2 data
```

	Date	Open	High	Low	Close	Volume
0	2010-11-04	291.00	344.90	291.00	342.55	479716245
1	2010-11-05	343.00	356.50	343.00	349.85	31927173
2	2010-11-08	351.80	355.90	329.50	330.75	46932779
3	2010-11-09	330.15	333.40	325.00	326.05	23741956
4	2010-11-10	325.40	327.80	320.05	322.80	21057129
...
2593	2021-04-26	127.75	128.00	126.50	126.95	4967884
2594	2021-04-27	127.00	127.90	126.60	127.50	3982954
2595	2021-04-28	128.00	129.45	127.50	128.50	6206074
2596	2021-04-29	129.75	130.05	127.65	128.05	8345584
2597	2021-04-30	127.40	134.60	127.00	133.05	27396950

2598 rows × 6 columns

```
1 from sklearn.model_selection import train_test_split
2 y = data['Close']
3 df = data
4 X = df.drop(['Close'], axis = 1)
5 X
```

	Date	Open	High	Low	Volume
0	2010-11-04	291.00	344.90	291.00	479716245
1	2010-11-05	343.00	356.50	343.00	31927173
2	2010-11-08	351.80	355.90	329.50	46932779
3	2010-11-09	330.15	333.40	325.00	23741956
4	2010-11-10	325.40	327.80	320.05	21057129
...
2593	2021-04-26	127.75	128.00	126.50	4967884
2594	2021-04-27	127.00	127.90	126.60	3982954
2595	2021-04-28	128.00	129.45	127.50	6206074
2596	2021-04-29	129.75	130.05	127.65	8345584
2597	2021-04-30	127.40	134.60	127.00	27396950

2598 rows × 5 columns

```
1 from sklearn.model_selection import train_test_split
2
3 X = []
4 Y = []
5 window_size=100
6 for i in range(1 , len(data) - window_size -1 , 1):
7     first = data.iloc[i,2]
8     temp = []
9     temp2 = []
10    for j in range(window_size):
11        temp.append((data.iloc[i + j, 2] - first) / first)
12        temp2.append((data.iloc[i + window_size, 2] - first) / first)
13    X.append(np.array(temp).reshape(100, 1))
14    Y.append(np.array(temp2).reshape(1, 1))
15
16 x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.2, shuffle=False)
17
18 train_X = np.array(x_train)
19 test_X = np.array(x_test)
20 train_Y = np.array(y_train)
21 test_Y = np.array(y_test)
22
23 train_X = train_X.reshape(train_X.shape[0],1,100,1)
24 test_X = test_X.reshape(test_X.shape[0],1,100,1)
25
26 print(len(train_X))
27 print(len(test_X))
```

3. Apply the LSTM architecture to the above dataset.

```

1 import tensorflow as tf
2 from tensorflow.keras.layers import Conv1D, LSTM, Dense, Dropout, Bidirectional, TimeDistributed
3 from tensorflow.keras.layers import MaxPooling1D, Flatten
4 from tensorflow.keras.regularizers import L1, L2
5 from tensorflow.keras.metrics import Accuracy
6 from tensorflow.keras.metrics import RootMeanSquaredError
7 from tensorflow.keras.optimizers import Adam
8
9 def lstm_model(lr, af):
10     model = tf.keras.Sequential()
11
12     # CNN layers
13     model.add(TimeDistributed(Conv1D(64, kernel_size=3, activation='relu', input_shape=(None, 100, 1))))
14     model.add(TimeDistributed(MaxPooling1D(2)))
15     model.add(TimeDistributed(Conv1D(128, kernel_size=3, activation='relu')))
16     model.add(TimeDistributed(MaxPooling1D(2)))
17     model.add(TimeDistributed(Conv1D(64, kernel_size=3, activation='relu')))
18     model.add(TimeDistributed(MaxPooling1D(2)))
19     model.add(TimeDistributed(Flatten()))
20     # model.add(Dense(5, kernel_regularizer=L2(0.01)))
21
22     # LSTM layers
23     model.add(Bidirectional(LSTM(100, activation = af, return_sequences=True)))
24     model.add(Dropout(0.5))
25     model.add(Bidirectional(LSTM(100, activation = af, return_sequences=False)))
26     model.add(Dropout(0.5))
27
28     #Final layers
29     model.add(Dense(1, activation='linear'))
30     model.compile(optimizer=Adam(learning_rate=lr), loss='mse', metrics=['mse', 'mae'])
31
32     history = model.fit(train_X, train_Y, validation_data=(test_X, test_Y), epochs=10, batch_size=40, verbose=1, shuffle =False)
33     return model

```

4 . Demonstrate the results with various learning rates.

5. Demonstrate the results with various Activation functions.

```

1 lr = [0.0001, 0.001, 0.01, 0.1]
2 af = ['relu', 'tanh', 'linear', 'sigmoid']

```

```

1 print('Model 1 with learning rate of 0.0001')
2 model1 = lstm_model(lr[0], af[1])
3 print('Model 2 with learning rate of 0.001\n')
4 model2 = lstm_model(lr[1], af[1])
5 print('Model 3 with learning rate of 0.01\n')
6 model3 = lstm_model(lr[2], af[1])
7 print('Model 4 with learning rate of 0.1\n')
8 model4 = lstm_model(lr[3], af[1])

```

Model 1 with learning rate of 0.0001

Epoch 1/10

50/50 [=====] - 13s 61ms/step - loss: 0.0208 - mse: 0.0208 - mae: 0.1118 - val_loss: 0.0317 - val_mse: 0.0317 - v

Epoch 2/10

50/50 [=====] - 1s 30ms/step - loss: 0.0178 - mse: 0.0178 - mae: 0.1014 - val_loss: 0.0193 - val_mse: 0.0193 - v

Epoch 3/10

50/50 [=====] - 1s 30ms/step - loss: 0.0115 - mse: 0.0115 - mae: 0.0783 - val_loss: 0.0108 - val_mse: 0.0108 - v

Epoch 4/10

50/50 [=====] - 2s 32ms/step - loss: 0.0083 - mse: 0.0083 - mae: 0.0683 - val_loss: 0.0096 - val_mse: 0.0096 - v

Epoch 5/10

50/50 [=====] - 1s 28ms/step - loss: 0.0071 - mse: 0.0071 - mae: 0.0635 - val_loss: 0.0085 - val_mse: 0.0085 - v

Epoch 6/10

50/50 [=====] - 1s 27ms/step - loss: 0.0060 - mse: 0.0060 - mae: 0.0583 - val_loss: 0.0076 - val_mse: 0.0076 - v

Epoch 7/10

50/50 [=====] - 1s 28ms/step - loss: 0.0052 - mse: 0.0052 - mae: 0.0542 - val_loss: 0.0068 - val_mse: 0.0068 - v

Epoch 8/10

50/50 [=====] - 1s 29ms/step - loss: 0.0050 - mse: 0.0050 - mae: 0.0525 - val_loss: 0.0063 - val_mse: 0.0063 - v

Epoch 9/10

50/50 [=====] - 1s 28ms/step - loss: 0.0046 - mse: 0.0046 - mae: 0.0508 - val_loss: 0.0061 - val_mse: 0.0061 - v

Epoch 10/10

50/50 [=====] - 1s 28ms/step - loss: 0.0044 - mse: 0.0044 - mae: 0.0494 - val_loss: 0.0059 - val_mse: 0.0059 - v

Model 2 with learning rate of 0.001

Epoch 1/10

50/50 [=====] - 13s 62ms/step - loss: 0.0163 - mse: 0.0163 - mae: 0.0941 - val_loss: 0.0099 - val_mse: 0.0099 - v

Epoch 2/10

50/50 [=====] - 1s 27ms/step - loss: 0.0066 - mse: 0.0066 - mae: 0.0617 - val_loss: 0.0074 - val_mse: 0.0074 - v

Epoch 3/10

50/50 [=====] - 1s 29ms/step - loss: 0.0069 - mse: 0.0069 - mae: 0.0605 - val_loss: 0.0077 - val_mse: 0.0077 - v

Epoch 4/10

50/50 [=====] - 1s 29ms/step - loss: 0.0057 - mse: 0.0057 - mae: 0.0592 - val_loss: 0.0111 - val_mse: 0.0111 - v

```
50/50 [=====] - 1s 29ms/step - loss: 0.0037 - mse: 0.0037 - mae: 0.0532 - val_loss: 0.0111 - val_mse: 0.0111 - v
Epoch 5/10
50/50 [=====] - 2s 32ms/step - loss: 0.0066 - mse: 0.0066 - mae: 0.0621 - val_loss: 0.0101 - val_mse: 0.0101 - v
Epoch 6/10
50/50 [=====] - 1s 30ms/step - loss: 0.0047 - mse: 0.0047 - mae: 0.0525 - val_loss: 0.0055 - val_mse: 0.0055 - v
Epoch 7/10
50/50 [=====] - 1s 29ms/step - loss: 0.0042 - mse: 0.0042 - mae: 0.0488 - val_loss: 0.0063 - val_mse: 0.0063 - v
Epoch 8/10
50/50 [=====] - 1s 27ms/step - loss: 0.0039 - mse: 0.0039 - mae: 0.0475 - val_loss: 0.0068 - val_mse: 0.0068 - v
Epoch 9/10
50/50 [=====] - 1s 29ms/step - loss: 0.0038 - mse: 0.0038 - mae: 0.0473 - val_loss: 0.0067 - val_mse: 0.0067 - v
Epoch 10/10
50/50 [=====] - 1s 28ms/step - loss: 0.0041 - mse: 0.0041 - mae: 0.0481 - val_loss: 0.0064 - val_mse: 0.0064 - v
Model 3 with learning rate of 0.01
```

```
Epoch 1/10
50/50 [=====] - 13s 61ms/step - loss: 0.0841 - mse: 0.0841 - mae: 0.1593 - val_loss: 0.0269 - val_mse: 0.0269 - v
Epoch 2/10
50/50 [=====] - 1s 28ms/step - loss: 0.0244 - mse: 0.0244 - mae: 0.1204 - val_loss: 0.0276 - val_mse: 0.0276 - v
Epoch 3/10
50/50 [=====] - 1s 28ms/step - loss: 0.0238 - mse: 0.0238 - mae: 0.1190 - val_loss: 0.0278 - val_mse: 0.0278 - v
Epoch 4/10
50/50 [=====] - 1s 27ms/step - loss: 0.0235 - mse: 0.0235 - mae: 0.1183 - val_loss: 0.0285 - val_mse: 0.0285 - v
Epoch 5/10
50/50 [=====] - 1s 27ms/step - loss: 0.0232 - mse: 0.0232 - mae: 0.1175 - val_loss: 0.0291 - val_mse: 0.0291 - v
Epoch 6/10
50/50 [=====] - 1s 29ms/step - loss: 0.0230 - mse: 0.0230 - mae: 0.1169 - val_loss: 0.0297 - val_mse: 0.0297 - v
Epoch 7/10
```

```
1 print('Model 5 with activation function - ReLU')
2 model5 = lstm_model(lr[1], af[0])
3 print('\nModel 6 with activation function - Tanh')
4 model6 = lstm_model(lr[1], af[1])
5 print('\nModel 7 with activation function - linear')
6 model7 = lstm_model(lr[1], af[2])
7 print('\nModel 8 with activation function - sigmoid')
8 model8 = lstm_model(lr[1], af[3])
```

```
Model 5 with activation function - ReLU
Epoch 1/10
50/50 [=====] - 9s 44ms/step - loss: 0.0202 - mse: 0.0202 - mae: 0.1099 - val_loss: 0.0198 - val_mse: 0.0198 - v
Epoch 2/10
50/50 [=====] - 1s 28ms/step - loss: 0.0103 - mse: 0.0103 - mae: 0.0765 - val_loss: 0.0136 - val_mse: 0.0136 - v
Epoch 3/10
50/50 [=====] - 1s 26ms/step - loss: 0.0070 - mse: 0.0070 - mae: 0.0645 - val_loss: 0.0130 - val_mse: 0.0130 - v
Epoch 4/10
50/50 [=====] - 1s 26ms/step - loss: 0.0070 - mse: 0.0070 - mae: 0.0628 - val_loss: 0.0090 - val_mse: 0.0090 - v
Epoch 5/10
50/50 [=====] - 1s 26ms/step - loss: 0.0049 - mse: 0.0049 - mae: 0.0541 - val_loss: 0.0084 - val_mse: 0.0084 - v
Epoch 6/10
50/50 [=====] - 1s 25ms/step - loss: 0.0048 - mse: 0.0048 - mae: 0.0532 - val_loss: 0.0081 - val_mse: 0.0081 - v
Epoch 7/10
50/50 [=====] - 1s 27ms/step - loss: 0.0049 - mse: 0.0049 - mae: 0.0530 - val_loss: 0.0091 - val_mse: 0.0091 - v
Epoch 8/10
50/50 [=====] - 1s 28ms/step - loss: 0.0042 - mse: 0.0042 - mae: 0.0501 - val_loss: 0.0075 - val_mse: 0.0075 - v
Epoch 9/10
50/50 [=====] - 1s 27ms/step - loss: 0.0047 - mse: 0.0047 - mae: 0.0521 - val_loss: 0.0064 - val_mse: 0.0064 - v
Epoch 10/10
50/50 [=====] - 1s 27ms/step - loss: 0.0037 - mse: 0.0037 - mae: 0.0467 - val_loss: 0.0067 - val_mse: 0.0067 - v
```

```
Model 6 with activation function - Tanh
Epoch 1/10
50/50 [=====] - 13s 64ms/step - loss: 0.0170 - mse: 0.0170 - mae: 0.0965 - val_loss: 0.0111 - val_mse: 0.0111 - v
Epoch 2/10
50/50 [=====] - 1s 27ms/step - loss: 0.0069 - mse: 0.0069 - mae: 0.0625 - val_loss: 0.0070 - val_mse: 0.0070 - v
Epoch 3/10
50/50 [=====] - 1s 26ms/step - loss: 0.0062 - mse: 0.0062 - mae: 0.0577 - val_loss: 0.0079 - val_mse: 0.0079 - v
Epoch 4/10
50/50 [=====] - 1s 27ms/step - loss: 0.0062 - mse: 0.0062 - mae: 0.0610 - val_loss: 0.0091 - val_mse: 0.0091 - v
Epoch 5/10
50/50 [=====] - 1s 26ms/step - loss: 0.0068 - mse: 0.0068 - mae: 0.0619 - val_loss: 0.0051 - val_mse: 0.0051 - v
Epoch 6/10
50/50 [=====] - 1s 26ms/step - loss: 0.0044 - mse: 0.0044 - mae: 0.0510 - val_loss: 0.0055 - val_mse: 0.0055 - v
Epoch 7/10
50/50 [=====] - 1s 26ms/step - loss: 0.0038 - mse: 0.0038 - mae: 0.0475 - val_loss: 0.0052 - val_mse: 0.0052 - v
Epoch 8/10
50/50 [=====] - 1s 26ms/step - loss: 0.0041 - mse: 0.0041 - mae: 0.0488 - val_loss: 0.0055 - val_mse: 0.0055 - v
Epoch 9/10
50/50 [=====] - 1s 27ms/step - loss: 0.0035 - mse: 0.0035 - mae: 0.0453 - val_loss: 0.0060 - val_mse: 0.0060 - v
Epoch 10/10
50/50 [=====] - 1s 26ms/step - loss: 0.0036 - mse: 0.0036 - mae: 0.0460 - val_loss: 0.0058 - val_mse: 0.0058 - v
```

```
Model 7 with activation function - linear
Epoch 1/10
50/50 [=====] - 10s 57ms/step - loss: 0.0170 - mse: 0.0170 - mae: 0.0966 - val_loss: 0.0104 - val_mse: 0.0104 - v
Epoch 2/10
50/50 [=====] - 1s 29ms/step - loss: 0.0071 - mse: 0.0071 - mae: 0.0620 - val_loss: 0.0073 - val_mse: 0.0073 - v
Epoch 3/10
50/50 [=====] - 2s 32ms/step - loss: 0.0057 - mse: 0.0057 - mae: 0.0558 - val_loss: 0.0065 - val_mse: 0.0065 - v
Epoch 4/10
50/50 [=====] - 1s 28ms/step - loss: 0.0051 - mse: 0.0051 - mae: 0.0563 - val_loss: 0.0096 - val_mse: 0.0096 - v
Epoch 5/10
```

Epoch 5/10

50/50 [=====] - 1s 27ms/step - loss: 0.0070 - mse: 0.0070 - mae: 0.0633 - val_loss: 0.0070 - val_mse: 0.0070 - val_mae: 0.0633

Epoch 6/10

50/50 [=====] - 1s 26ms/step - loss: 0.0059 - mse: 0.0059 - mae: 0.0592 - val_loss: 0.0072 - val_mse: 0.0072 - val_mae: 0.0592

Epoch 7/10

6. Demonstrate the results with all possible evaluation criteria.

```
1 res_lr = []
2 res_lr.append(model1.evaluate(test_X, test_Y))
3 res_lr.append(model2.evaluate(test_X, test_Y))
4 res_lr.append(model3.evaluate(test_X, test_Y))
5 res_lr.append(model4.evaluate(test_X, test_Y))
6 modelNums = ['\n-----Model1-----', '\n-----Model2-----', '\n-----Model3-----', '\n-----Model4-----']
7 para = ['Test Loss : ', 'MSE : ', 'MAE : ']
8 for i in range(4):
9     print(modelNums[i])
10    for j in range(3):
11        print(para[j], res_lr[i][j])
```

16/16 [=====] - 0s 8ms/step - loss: 0.0059 - mse: 0.0059 - mae: 0.0622

16/16 [=====] - 0s 9ms/step - loss: 0.0064 - mse: 0.0064 - mae: 0.0648

16/16 [=====] - 0s 8ms/step - loss: 0.0302 - mse: 0.0302 - mae: 0.1453

16/16 [=====] - 0s 7ms/step - loss: 0.0887 - mse: 0.0887 - mae: 0.2605

-----Model1-----

Test Loss : 0.00589575432240963

MSE : 0.00589575432240963

MAE : 0.06219031661748886

-----Model2-----

Test Loss : 0.006415408104658127

MSE : 0.006415408104658127

MAE : 0.06482697278261185

-----Model3-----

Test Loss : 0.030163953080773354

MSE : 0.030163953080773354

MAE : 0.14534984529018402

-----Model4-----

Test Loss : 0.08865439891815186

MSE : 0.08865439891815186

MAE : 0.26054394245147705

```
1 res_af = []
2 res_af.append(model5.evaluate(test_X, test_Y))
3 res_af.append(model6.evaluate(test_X, test_Y))
4 res_af.append(model7.evaluate(test_X, test_Y))
5 res_af.append(model8.evaluate(test_X, test_Y))
6 modelNums = ['\n-----Model5(ReLU)-----', '\n-----Model6(Tanh)-----', '\n-----Model7(Linear)-----', '\n-----Model8(Sigmoid)-----']
7 para = ['Test Loss : ', 'MSE : ', 'MAE : ']
8 for i in range(4):
9     print(modelNums[i])
10    for j in range(3):
11        print(para[j], res_af[i][j])
```

16/16 [=====] - 0s 7ms/step - loss: 0.0067 - mse: 0.0067 - mae: 0.0670

16/16 [=====] - 0s 8ms/step - loss: 0.0058 - mse: 0.0058 - mae: 0.0610

16/16 [=====] - 0s 7ms/step - loss: 0.0074 - mse: 0.0074 - mae: 0.0703

16/16 [=====] - 0s 7ms/step - loss: 0.0119 - mse: 0.0119 - mae: 0.0888

-----Model5(ReLU)-----

Test Loss : 0.006718510761857033

MSE : 0.006718510761857033

MAE : 0.06697747856378555

-----Model6(Tanh)-----

Test Loss : 0.005775176454335451

MSE : 0.005775176454335451

MAE : 0.061000458896160126

-----Model7(Linear)-----

Test Loss : 0.0073815868236124516

MSE : 0.0073815868236124516

MAE : 0.07033152133226395

-----Model8(Sigmoid)-----

Test Loss : 0.011858091689646244

MSE : 0.011858091689646244

MAE : 0.08875854313373566

8. Predict the stock market value based on your implemented model.

```
1 from sklearn.metrics import explained_variance_score, mean_poisson_deviance, mean_gamma_deviance
2 from sklearn.metrics import r2_score
```

```

3 from sklearn.metrics import max_error
4
5 yhat_probs = model2.predict(test_X, verbose=0)
6 yhat_probs = yhat_probs[:, 0]
7 print('Model with learning rate = 0.001 and activation function - Tanh\n')
8 var = explained_variance_score(test_Y.reshape(-1,1), yhat_probs)
9 print('Variance: %f' % var)
10
11 r2 = r2_score(test_Y.reshape(-1,1), yhat_probs)
12 print('R2 Score: %f' % var)
13
14 var2 = max_error(test_Y.reshape(-1,1), yhat_probs)
15 print('Max Error: %f' % var2)

```

Model with learning rate = 0.001 and activation function - Tanh

Variance: 0.818686

R2 Score: 0.818686

Max Error: 0.246054

9. Demonstrate the learning results in all possible visualizing ways.

```

1 predicted = model6.predict(test_X)
2 test_label = test_Y.reshape(-1,1)
3 predicted = np.array(predicted[:,0]).reshape(-1,1)
4 len_t = len(train_X)
5 for j in range(len_t, len_t + len(test_X)):
6     temp = data.iloc[j,3]
7     test_label[j - len_t] = test_label[j - len_t] * temp + temp
8     predicted[j - len_t] = predicted[j - len_t] * temp + temp
9 plt.plot(predicted, color = 'green', label = 'Predicted Stock Price')
10 plt.plot(test_label, color = 'red', label = 'Real Stock Price')
11 plt.title(' Stock Price Prediction')
12 plt.xlabel('Time')
13 plt.ylabel(' Stock Price')
14 plt.legend()
15 plt.show()

```

16/16 [=====] - 2s 7ms/step

